

Kernel Method and Support Vector Machines

Nguyen Duc Dung, Ph.D.
Institute of Information Technology, VAST

Outline

■ Reference

- Books, papers, slides, software

■ Support vector machines (SVMs)

- The maximum-margin hyper-plane
- Kernel method

■ Implementation

- Approaches
- Sequential minimal optimization (SMO)

■ Open problems

■ SVMs Approximation

- Reduced set method

Reference

■ Book

- Cristianini, N., Shawe-Taylor, J., *An Introduction to Support Vector Machines*, Cambridge University Press, (2000).
<http://www.support-vector.net/index.html>
- Bernhard Schölkopf and Alex Smola. **Learning with Kernels.** MIT Press, Cambridge, MA, 2002.

■ Paper

- C. J. C. Burges. **A Tutorial on Support Vector Machines for Pattern Recognition.** *Knowledge Discovery and Data Mining*, 2(2), 1998.

■ Slide

- N. Cristianini. **ICML'01 tutorial**, 2001.


■ Software

- LibSVM (NTU), SVM^{light} (joachims.org)

■ Online resource


- <http://www.kernel-machines.org/>

Classification Problem




How would we classify this data set?

Linear Classifiers



There are many lines that can be linear classifiers.
Which one is the better classifier?

Margin of a Linear Classifier



- **Margin:** The **minimum distance** to the training data
- **SVM solution:** The linear classifier with **maximum margin**

Margin

of a Linear Function $f(x) = w \cdot x + b$

- Functional margin

$$\hat{\rho}_f(\mathbf{x}_i, y_i) = y_i(\mathbf{w} \cdot \mathbf{x}_i + b)$$


- Geometric margin

$$\rho_f(\mathbf{x}_i, y_i) = \frac{\hat{\rho}_f(\mathbf{x}_i, y_i)}{\|\mathbf{w}\|}$$

- Margin $\rho_f = \min_{i=1 \dots l} \rho_f(\mathbf{x}_i, y_i)$

- SVM solution

$$f^* = \arg \max_f \rho_f$$



A Bound on Expected Risk of a Linear Classifier $f = \text{sign}(w \cdot x)$

With a probability at least $(1 - \delta)$, $\delta \in (0, 1)$


$$R[f] \leq R_{\text{emp}}[f] + \sqrt{\frac{c}{l} \left(\frac{R^2 \Lambda^2}{\rho_f^2} \ln^2 l + \ln \frac{1}{\delta} \right)}$$

where R_{emp} is training error, l is training size, ρ_f is the margin, $\|w\| \leq \Lambda$, $\|x\| \leq R$, c is a constant

Larger margin, smaller bound

Finding the Maximum-Margin Classifier


$$f^* = \arg \max_f \rho_f$$




minimize $\frac{1}{2} \|\mathbf{w}\|^2,$
subject to $y_i(\mathbf{w} \cdot \mathbf{x}_i + b) \geq 1, i = 1, \dots, l$

Minimize normal vector

Constrain functional margin ≥ 1




Soft and Hard Margin



$$\begin{aligned} \min_{w,b} \quad & \frac{1}{2} \|w\|^2 \\ \text{s.t.} \quad & y_i(w \cdot x_i + b) \geq 1, i = 1, \dots, l \end{aligned}$$

Hard (maximum) margin



$$\begin{aligned} \min_{w,b} \quad & \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \xi_i^p \\ \text{s.t.} \quad & y_i(w \cdot x_i + b) \geq 1 - \xi_i, \\ & \xi_i \geq 0, i = 1, \dots, l \end{aligned}$$

Soft (maximum) margin

Lagrangian Optimization

Definition 1 *Given an optimization problem with convex domain $\Omega \subseteq \mathbb{R}^d$*

$$\text{minimize} \quad f(w), w \in \Omega \quad (2.20)$$

$$\text{subject to} \quad g_i(w) \leq 0, i = 1, \dots, k \quad (2.21)$$

$$h_i(w) = 0, i = 1, \dots, m \quad (2.22)$$

The generalized Lagrangian function is defined as

$$L(w, \alpha, \beta) = f(w) + \sum_{i=1}^k \alpha_i g_i(w) + \sum_{i=1}^m \beta_i h_i(w) \quad (2.23)$$

Definition 2 *The Lagrangian dual problem of the primal problem is the following problem*

$$\text{maximize} \quad \theta(\alpha, \beta) \quad (2.24)$$

$$\text{subject to} \quad \alpha \geq 0 \quad (2.25)$$

where $\theta(\alpha, \beta) = \inf_{w \in \Omega} L(w, \alpha, \beta)$

Kuhn-Tucker Theorem

Theorem 3 (Kuhn-Tucker) *Given an optimization problem with convex domain $\Omega \subseteq \mathbb{R}^d$*

$$\text{minimize} \quad f(\mathbf{w}), \mathbf{w} \in \Omega \quad (2.27)$$

$$\text{subject to} \quad g_i(\mathbf{w}) \leq 0, i = 1, \dots, k \quad (2.28)$$

$$h_i(\mathbf{w}) = 0, i = 1, \dots, m \quad (2.29)$$

with $f \in C^1$ convex and g_i, h_i affine, necessary and efficient conditions for a normal point \mathbf{w}^ to be optimum are the existence of α^* and β^* such that*

$$\frac{\partial L(\mathbf{w}^*, \alpha^*, \beta^*)}{\partial \mathbf{w}} = 0 \quad (2.30)$$

$$\frac{\partial L(\mathbf{w}^*, \alpha^*, \beta^*)}{\partial \beta} = 0 \quad (2.31)$$

$$\alpha_i^* g_i(\mathbf{w}^*) = 0, i = 1, \dots, k \quad (2.32)$$

$$g_i(\mathbf{w}^*) \leq 0, i = 1, \dots, k \quad (2.33)$$

$$\alpha_i \geq 0, i = 1, \dots, k \quad (2.34)$$

Optimization

$$\begin{aligned} \min_{w,b} \quad & \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \xi_i^p \\ \text{s.t.} \quad & y_i(wx_i + b) \geq 1 - \xi_i, \\ & \xi_i \geq 0, i = 1, \dots, l \end{aligned}$$

Primal problem

$$L(w, b, \xi, \alpha, \beta) = \frac{1}{2} w^2 + C \sum_{i=1}^l \xi_i - \sum_{i=1}^l \alpha_i (y_i(w \cdot x_i + b) - 1 + \xi_i) - \sum_{i=1}^l \beta_i \xi_i$$

$$\frac{\partial L(w, \alpha, \beta)}{\partial w} = w - \sum_{i=1}^l y_i \alpha_i x_i = 0$$

$$\mathbf{w} = \sum_{\alpha_i \neq 0} y_i \alpha_i x_i$$

$$\frac{\partial L(w, \alpha, \beta)}{\partial \xi_i} = C - \alpha_i - \beta_i = 0$$

$$\frac{\partial L(w, \alpha, \beta)}{\partial b} = \sum_{i=1}^l y_i \alpha_i = 0$$

Dual problem

$$\begin{aligned} \min_{\alpha_i} \quad & \frac{1}{2} \sum_{i,j=1}^l y_i y_j \alpha_i \alpha_j \langle x_i, x_j \rangle - \sum_{i=1}^l \alpha_i \\ \text{s.t.:} \quad & 0 \leq \alpha_i \leq C, i = 1, \dots, l, \\ & \sum_{i=1}^l y_i \alpha_i = 0. \end{aligned}$$

(Linear) Support Vector Machines

■ Training

$$\min_{\alpha_i} \frac{1}{2} \sum_{i,j=1}^l y_i y_j \alpha_i \alpha_j \langle x_i, x_j \rangle - \sum_{i=1}^l \alpha_i$$

$$\text{s.t.: } 0 \leq \alpha_i \leq C, i = 1, \dots, l,$$

$$\sum_{i=1}^l y_i \alpha_i = 0.$$

- Quadratic optimization
- l variables
- l^2 coefficients

■ Testing

$$f(x) = w \cdot x + b$$

- Norm of the hyperplane

$$\mathbf{w} = \sum_{\alpha_i \neq 0} y_i \alpha_i x_i$$

- $(x_i, \alpha_i), \alpha_i \neq 0$ – **support vector**


Kernel Method

■ Problem

- ❑ Most datasets are **linearly non-separable**


■ Solution

- ❑ Map input data into a *higher dimensional* feature space
- ❑ Find the optimal hyperplane in feature space



Hyperplane in Feature Space

- ❖ VC-dimension of a class of functions: the maximum number of points that can be shattered
- ❖ VC-dimension of linear functions in R^d is $d+1$
- ❖ Dimension of feature space is high
- **Linear functions** in feature space has high VC-dimension, or **high capacity**



VC Dimension: Example



Gaussian RBF SVMs of sufficiently small width can classify an arbitrary large number of training points correctly, and thus ***have infinite VC dimension***

Linear SVMs

■ Training

$$\min_{\alpha_i} \frac{1}{2} \sum_{i,j=1}^l y_i y_j \alpha_i \alpha_j \langle x_i, x_j \rangle - \sum_{i=1}^l \alpha_i$$

s.t.: $0 \leq \alpha_i \leq C, i = 1, \dots, l,$

$$\sum_{i=1}^l y_i \alpha_i = 0.$$

- Quadratic optimization
- l variables
- l^2 coefficients

■ Testing

$$f(x) = \text{sign} \left(\sum_{\alpha_i \neq 0} y_i \alpha_i \langle x, x_i \rangle + b \right)$$

- Norm of the hyperplane

$$w = \sum_{\alpha_i \neq 0} y_i \alpha_i x_i$$

- $(x_i, \alpha_i), \alpha_i \neq 0$ – support vector

SVMs work with **pairs** of data (dot product), not sample

Non-linear SVMs

- **Kernel**: to calculate dot product between two vectors in feature space $K(x,y) = \langle \Phi(x), \Phi(y) \rangle$

■ Training

$$\begin{aligned} \min_{\alpha_i} \quad & \frac{1}{2} \sum_{i,j=1}^l y_i y_j \alpha_i \alpha_j K(x_i, x_j) - \sum_{i=1}^l \alpha_i \\ \text{s.t.:} \quad & 0 \leq \alpha_i \leq C, i = 1, \dots, l, \\ & \sum_{i=1}^l y_i \alpha_i = 0. \end{aligned}$$

■ Testing

$$f(x) = \text{sign} \left(\sum_{\alpha_i \neq 0} y_i \alpha_i K(x, x_i) + b \right)$$

Norm of the hyperplane

$$\Psi = \sum_{\alpha_i \neq 0} y_i \alpha_i \Phi(x_i)$$

The maximal margin algorithm works indirectly in feature space via kernel, or Φ is not known explicitly

Kernel

- Linear: $K(x,y) = \langle x,y \rangle$
- Gaussian: $K(x,y) = \exp(-\gamma\|x-y\|^2)$
 - Dimension of feature space: *infinite*
- Polynomial: $K(x,y) = \langle x,y \rangle^p$
 - Dimension of feature space: $\binom{d+p-1}{p}$, where d – input space dimension

Theorem 4 (Mercer) *To guarantee that a continuous symmetric function $K(u,v)$ in $L_2(C)$ has an expansion*

$$K(u,v) = \sum_{i=1}^{\infty} a_k z_k(u) z_k(v) \quad (2.53)$$

with positive coefficients $a_k > 0$ (i.e., $K(u,v)$ describes an inner product in some feature space), it is necessary and sufficient that the condition


$$\int_C \int_C K(u,v) g(u) g(v) du dv \geq 0 \quad (2.54)$$

is valid for all $g \in L_2(C)$ (C being a compact subset of \mathbb{R}^d)

Support Vector Learning

■ Task

- Given a set of labeled data
 $T = \{(x_i, y_i)\}_{i=1, \dots, l} \subset R^d \times \{-1, +1\}$
- Find the decision function



■ Training

**Time: $O(l^3)$,
Memory: $O(l^2)$**

$$\begin{aligned} \min_{\alpha_i} \quad & \frac{1}{2} \sum_{i,j=1}^l y_i y_j \alpha_i \alpha_j K(x_i, x_j) - \sum_{i=1}^l \alpha_i \\ \text{s.t.:} \quad & 0 \leq \alpha_i \leq C, i = 1, \dots, l, \\ & \sum_{i=1}^l y_i \alpha_i = 0. \end{aligned}$$

■ Testing


Time: $O(Ns)$

$$f(x) = \text{sign} \left(\sum_{\alpha_i \neq 0} y_i \alpha_i K(x, x_i) + b \right)$$

MNIST Data: SVM vs. Other

- Data
 - 60,000/10,000 training/testing
- Performance

Method	Testing error (%)
linear classifier (1-layer NN)	12.0
K-nearest-neighbors	5.0
40 PCA + quadratic classifier	3.3
<i>SVM, Gaussian Kernel</i>	<i>1.4</i>
2-layer NN, 300 hidden units, mean square error	4.7
Convolutional net LeNet-4	1.1



Hand written data

(Source: <http://yann.lecun.com/>)

SVM: Probability Output

- SVM solution

$$f(x) = \sum_{\alpha_i \neq 0} y_i \alpha_i K(x_i, x) + b$$

- Probability estimation

$$p(y = +1 | x) \approx \frac{1}{1 + e^{Af(x)+B}}$$

- Maximum likelihood approach

$$(A, B) = \arg \min_{a, b} F(a, b) = - \sum_{i=1}^l (t_i \log(p_i) + (1 - t_i) \log(1 - p_i))$$

where $p_i = p(y = +1 | x_i) \approx \frac{1}{1 + e^{af(x)+b}},$

$$t_i = \begin{cases} \frac{N_+ + 1}{N_+ + 2} & \text{if } y_i = +1, \\ \frac{1}{N_- + 1} & \text{if } y_i = -1 \end{cases}, i = 1, \dots, l. (N_+ : \# \text{positive}, N_- : \# \text{negative})$$

Outline

- **Reference**

- Books, papers, slides, software

- **Support vector machines (SVMs)**

- The maximum-margin hyperplane
- Kernel method

- **Implementation**

- Approaches
- Sequential minimal optimization

- **Open problems**

SVM Training

Problem

$$\begin{aligned} \min_{\alpha_i} F(\alpha) &= \frac{1}{2} \sum_{i,j=1}^l y_i y_j \alpha_i \alpha_j K_{ij} - \sum_{i=1}^l \alpha_i \\ \text{s.t.: } 0 &\leq \alpha_i \leq C, i = 1, \dots, l, \\ \sum_{i=1}^l \alpha_i y_i &= 0 \end{aligned}$$

Quadratic programming (QP)

- Obj. function: quadratic w.r.t. α
- Number of variable: l
- Number of parameter: l^2
- Complexity
 - Time: $O(l^3)$ or $O(N_S^3 + N_S^2 l + N_S d l)$
 - Memory: $O(l^2)$
- Constraint: box, linear

Approach

■ Gradient method

- Modified gradient projection (Bottou et al., 94)

■ Divide-and-conquer

- Decomposition alg. (e.g. Osuna et al., 97, Joachims, 99)
- Sequential minimal optimization (SMO) (Plat, 99)

■ Parallelization

- Cascade SVM (Peter et al., 05)
- Parallel mixture of SVM (Collobert et al., 02)

■ Approximation

- Online and active learning (e. g. Bordes et al., 05)
- Core SVM (Tsang et al., 05, 07)


■ Combination of methods

Optimality

The Karush-Kuhn-Tucker (KKT) conditions

$$\begin{cases} y_i f(x_i) > 1 & \text{for } \alpha_i = 0, \\ y_i f(x_i) < 1 & \text{for } \alpha_i = C, \\ y_i f(x_i) < 1 & \text{for } 0 < \alpha_i < C, \end{cases}$$

where $f(x) = \sum_{i=1}^l y_i \alpha_i K(x, x_i) + b$



SMO Algorithm

- Initialize solution (zero)
- **While** (*!StoppingCondition*)
 - Select two vector $\{i,j\}$
 - Optimize on $\{i,j\}$
- **EndWhile**

SMO: Optimization

■ Problem

$$\min_{\alpha_i} F(\boldsymbol{\alpha}) = \frac{1}{2} \sum_{i,j=1}^l y_i y_j \alpha_i \alpha_j K_{ij} - \sum_{i=1}^l \alpha_i$$

$$\text{s.t.: } 0 \leq \alpha_i \leq C, i = 1, \dots, l,$$

$$\sum_{k=1}^l \alpha_k y_k = 0$$

$$\rightarrow \forall(i, j): y_i \alpha_i + y_j \alpha_j = \text{const}$$

$$\rightarrow \alpha_j = y_j (\text{const} - y_i \alpha_i)$$

■ Fixing all $\alpha_k, k \neq i, j$

$$F(\boldsymbol{\alpha}) = F(\alpha_i) = A\alpha_i^2 + B\alpha_i + C$$

■ Updating scheme (without the box constraint)

$$\alpha_i^{\text{new}} = \alpha_i^{\text{old}} + \frac{y_i (E_j^{\text{old}} - E_i^{\text{old}})}{2\kappa_{ij}},$$
$$\alpha_j^{\text{new}} = \alpha_j^{\text{old}} + \frac{y_j (E_i^{\text{old}} - E_j^{\text{old}})}{2\kappa_{ij}}.$$

$$E_i = \sum_{k=1}^l y_k \alpha_k K(x_k, x_i) - y_i, i = 1, \dots, l,$$

$$\kappa_{ij} = K_{ii} + K_{jj} - 2K_{ij}$$

Selection Heuristic and Stopping Condition

- Maximum violating pair

$$\begin{cases} i = \arg \max \{-E_k \mid k \in I_{up}\} \\ j = \arg \min \{-E_k \mid k \in I_{low}\} \end{cases}$$

- Maximum gain

$$\begin{cases} i = \arg \max \{-E_k \mid k \in I_{up}\} \\ j = \arg \max \{|\Delta F_{ik}| \mid k \in I_{low}, -E_k < -E_i\} \end{cases}$$

where $I_{up} = \{t \mid \alpha_t < C, y_t = +1 \text{ or } \alpha_t > 0, y_t = -1\}$

$I_{low} = \{t \mid \alpha_t < C, y_t = -1 \text{ or } \alpha_t > 0, y_t = +1\}$

- Stopping condition: $|E_i - E_j| < \varepsilon(10^{-3})$

Sequential Minimal Optimization

■ Training problem

$$\min_{\alpha_i} \frac{1}{2} \sum_{i,j=1}^l y_i y_j \alpha_i \alpha_j K(x_i, x_j) - \sum_{i=1}^l \alpha_i$$

$$\text{s.t.: } 0 \leq \alpha_i \leq C, i = 1, \dots, l,$$

$$\sum_{i=1}^l y_i \alpha_i = 0.$$

■ Functional margin

$$E_i = \sum_{k=1}^l y_k \alpha_k K(x_k, x_i) - y_i$$

■ Selection heuristic

$$i = \arg \max_k \{-E_k \mid k \in I_{up}(\alpha)\}$$

$$j = \arg \max_k \{|\Delta L_{ik}| \mid k \in I_{low}(\alpha), E_k < E_i\}$$

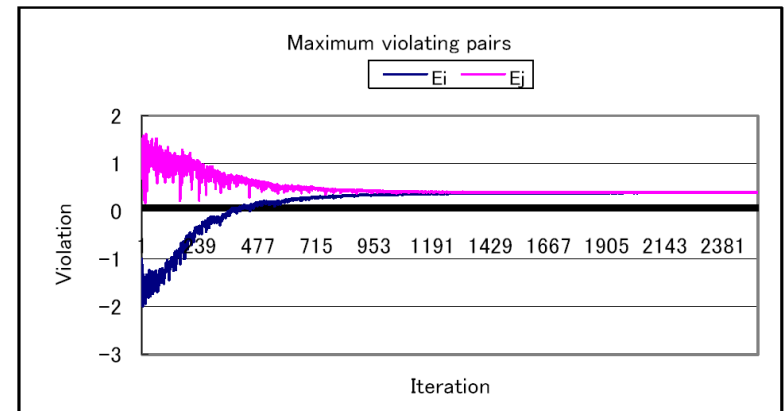
■ Updating scheme

$$\alpha_i^{new} = \alpha_i^{old} + \frac{y_i (E_j^{old} - E_i^{old})}{2\kappa_{ij}},$$

$$\alpha_j^{new} = \alpha_j^{old} + \frac{y_j (E_i^{old} - E_j^{old})}{2\kappa_{ij}}.$$


■ Stopping condition

$$|E_i - E_j| < \varepsilon$$



Support Vector Regression (1)

- Training data $S = \{(x_i, y_i)\}_{i=1, \dots, l} \subset \mathbf{R}^N \times \mathbf{R}$
- Linear regressor $y = f(x) = \mathbf{w} \cdot \mathbf{x} + b$
- ε -loss function



$$L^\varepsilon((x_i, y_i), f) = |y_i - f(x_i)|_\varepsilon = \max(0, |y_i - f(x_i)| - \varepsilon)$$

Support Vector Regression (2)

- Optimization: minimizing

$$\frac{1}{2} \|\mathbf{w}\|_2^2 + C \sum_{i=1}^l L^\epsilon((\mathbf{x}_i, y_i), f)$$

- Dual problem

$$\begin{aligned} \text{maximize} \quad & \sum_{i=1}^l (\alpha_i^- - \alpha_i^+) y_i - \epsilon \sum_{i=1}^l (\alpha_i^- + \alpha_i^+) \\ & - \frac{1}{2} \sum_{i,j=1}^l l(\alpha_i^- - \alpha_i^+) (\alpha_j^- - \alpha_j^+) \mathbf{x}_i \cdot \mathbf{x}_j, \\ \text{subject to} \quad & 0 \leq \alpha_i^+, \alpha_i^- \leq C, i = 1, \dots, l \\ & \sum_{i=1}^l (\alpha_i^- - \alpha_i^+) = 0, i = 1, \dots, l \end{aligned}$$

Open Problems

■ Model selection

- Kernel type
- Parameter setting

■ Speed and size

- Training: time $O(N_S^2 l)$,
space $O(N_S l)$
- Testing: $O(N_S)$

■ Multi-class application

- One-versus-rest
- One-versus-one

■ Categorical data

Open Problems

■ Model selection

- Kernel type
- Parameter setting


■ Speed and size

- Training: time $O(N_S^2 l)$, space $O(N_S l)$
- Testing: $O(N_S)$

■ Multi-class application

- One-versus-rest
- One-versus-one

■ Categorical data



Reduced Set Method

To replace original machine

$$y = \text{sign} \left(\sum_{i=1}^{N_s} \alpha_i K(x_i, x) + b \right) \quad (1)$$

By a simplified machine

$$y' = \text{sign} \left(\sum_{j=1}^{N_z} \beta_j K(z_j, x) + b \right) \quad (2)$$


with

$$N_z \ll N_s,$$

(1) and (2) are similar

$\{(x_i, \alpha_i)\}_{i=1, \dots, N_s}$ – original vectors

$\{(z_j, \beta_j)\}_{j=1, \dots, N_z}$ – reduced vectors



Performance

■ Dataset

Name	Dimension	# Class	# Training	#Testing
DNA	180	3	2,000	1,186
Satimage	36	6	4,435	2,000
Shuttle	9	7	43,500	14,500
USPS	256	10	7,291	2,007

■ Result

Data	DNA		Satimage		Shuttle		USPS	
% of SV	#SV	Acc. (%)	#SV	Acc. (%)	#SV	Acc. (%)	#SV	Acc. (%)
100%	843	95.62	1215	89.75	4191	99.03	1670	94.77
50%	422	95.62	608	89.75	2096	99.03	835	94.77
10%	84	95.53	122	89.45	419	99.03	167	94.67
5%	42	95.19	61	89.25	210	99.03	84	93.92
1%	8	95.03	12	78.00	42	99.04	45	89.59

High reduction rate, no change in predictive accuracy

Practice

- **Download the following SVM tools**
 - LibSVM: [LIBSVM -- A Library for Support Vector Machines \(ntu.edu.tw\)](http://ntu.edu.tw/~csie/libsvm/)
- **Select one classification problem**
 - From UCI repository
 - From LibSVM data sets
- **Report on progress in selected problem**
 - State-of-the-art works
 - SVM: model selection, performance